

nfsroot

overview and discussion points

Jim Garlick

garlick@llnl.gov

Livermore Computing

Lawrence Livermore National Laboratory

LLNL-PRES-615396



Presentation Topics

- overview
- diskless boot sequence walkthrough
- scalability questions for large clusters
- distributed network block device

overview

nfsroot scope

nfsroot makes a root fs image sharable via NFS

The main unique thing about nfsroot is it is contained in the diskless root image. There is no server component, therefore configuring DHCP, TFTP, NFS, and building images is not in scope.

nfsroot mainly consists of

- scripts to configure boot payloads, chrooted on server
- scripts to make RO root image appear RW, on client
- initrd tools [RHEL6: handled by dracut]

boot payloads

PXE payload

```
/boot/pxelinux.0  
/boot/pxelinux.cfg  
/boot/pxelinux.msg
```

Alternative OS payload

```
/boot/freedos.img  
/boot/memdisk  
/boot/memtest86+-4.10  
/boot/memtest86+ -> memtest86+-4.10
```

Linux payload

```
/boot/vmlinuz -> vmlinuz-2.6.32-220.23.1.1chaos.ch5.x86_64  
/boot/System.map -> System.map-2.6.32-220.23.1.1chaos.ch5.x86_64  
/boot/initramfs -> initramfs-2.6.32-220.23.1.1chaos.ch5.x86_64.img  
/boot/vmlinuz-2.6.32-220.23.1.1chaos.ch5.x86_64  
/boot/System.map-2.6.32-220.23.1.1chaos.ch5.x86_64  
/boot/initramfs-2.6.32-220.23.1.1chaos.ch5.x86_64.img
```

Kdump payload

```
/boot/initrd-2.6.32-220.23.1.1chaos.ch5.x86_64kdump.img ->  
initramfs-2.6.32-220.23.1.1chaos.ch5.x86_64.img
```



payload reconfig

nfsroot scripts

`/usr/sbin/configpxe`

`/usr/sbin/nfsroot-rebuild`

`/usr/sbin/nfsroot-setdefault`

`/usr/sbin/nfsroot-kdumplinks`

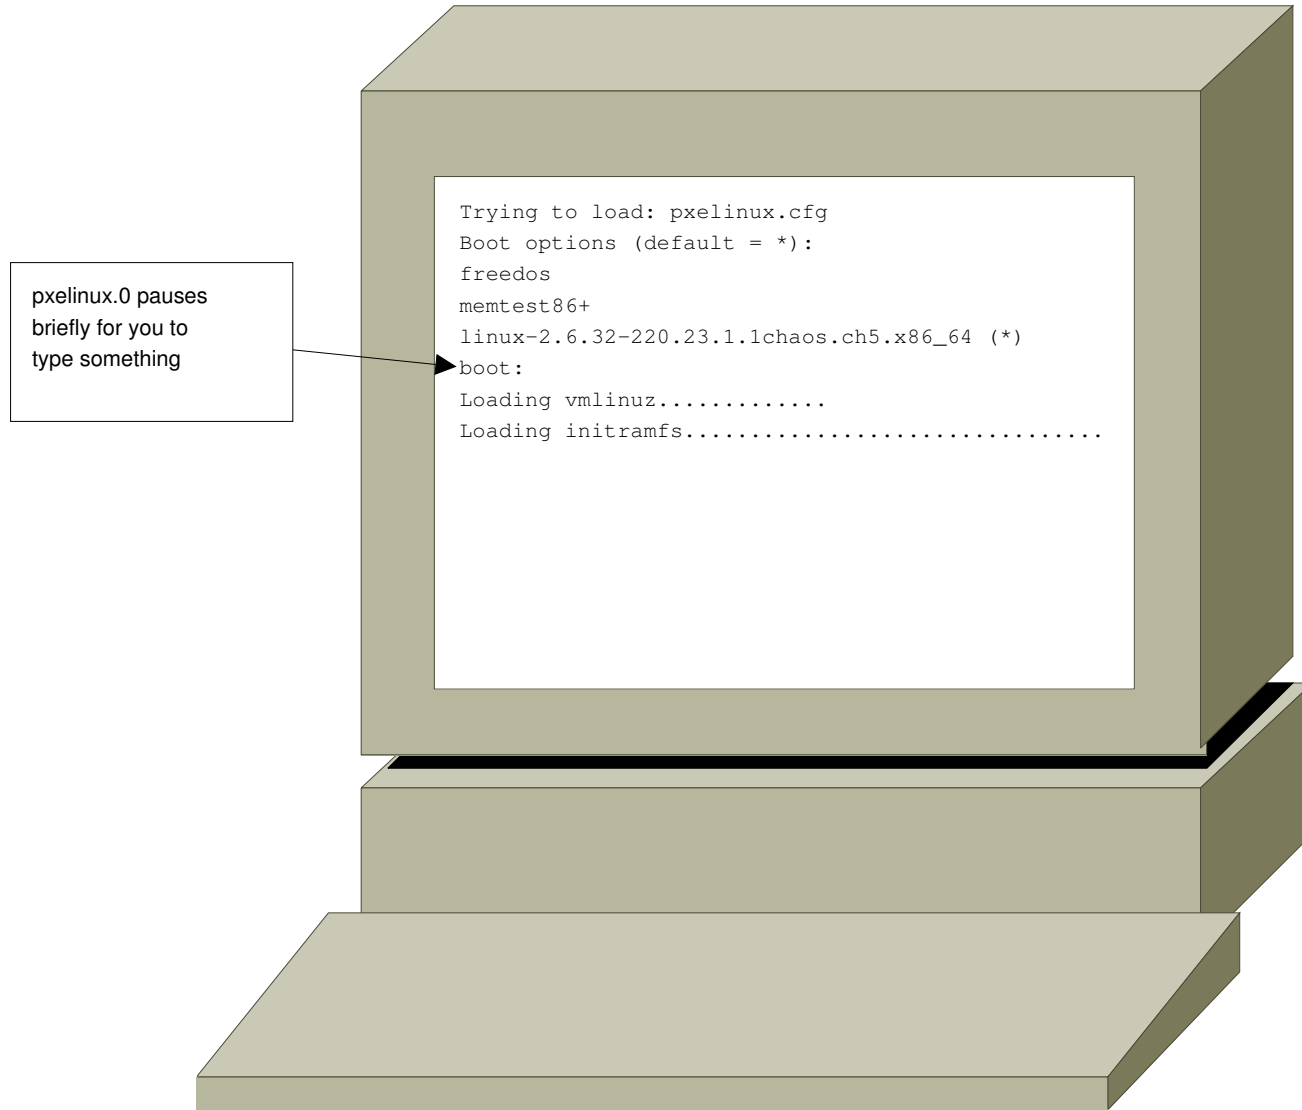
`/usr/sbin/nfsroot-memtestlinks`

called from grubby `/sbin/new-kernel-pkg`

`/etc/kernel/postinst.d/nfsroot-postinst`

`/etc/kernel/prerm.d/nfsroot-prerm`

pxelinux.0 boot prompt



making RO root appear RW

`/etc/rc.nfsroot` tries boot methods until one succeeds

- **unionfs** - overlay tmpfs
- **aufs** - overlay tmpfs
- **bind** - bind mount
tmpfs copies of dirs
- **bindnfs** - bind mount
NFS copies of dirs
- **rbind** - bind mount RO
dirs into tmpfs /
- **ram** - copy whole root
into tmpfs
- **none** - non-shared RW
root
- **kdump** - save kdump
to NFS and reboot
- **zram** - RO network
block device root with
zram overlay

nfsroot configuration

/etc/sysconfig/nfsroot configures client boot behavior

```
# methods are tried in this order
```

```
METHODS="kdump none aufs unionfs zram bind ram"
```

```
# tmpfs max size (if tmpfs used)
```

```
#TMPFSMAX=128m
```

```
# bind only: dirs to bind mount, subdirs to not copy
```

```
RAMDIRS="/etc /var /mnt /root"
```

```
#RAMDIRS_NOCOPY="/var/cache/yum /var/lib/rpm /var/lib/yum"
```

```
# kdump config
```

```
#KDUMP_DIR=disthost:/tftpboot/dumps
```

```
#KDUMP_DIR_MOUNTOPTS="nfsvers=3,rw,nolock"
```

```
#KDUMP_LEVEL=31
```

```
KDUMP_FAILSAFE=shell
```


diskless boot sequence walkthrough



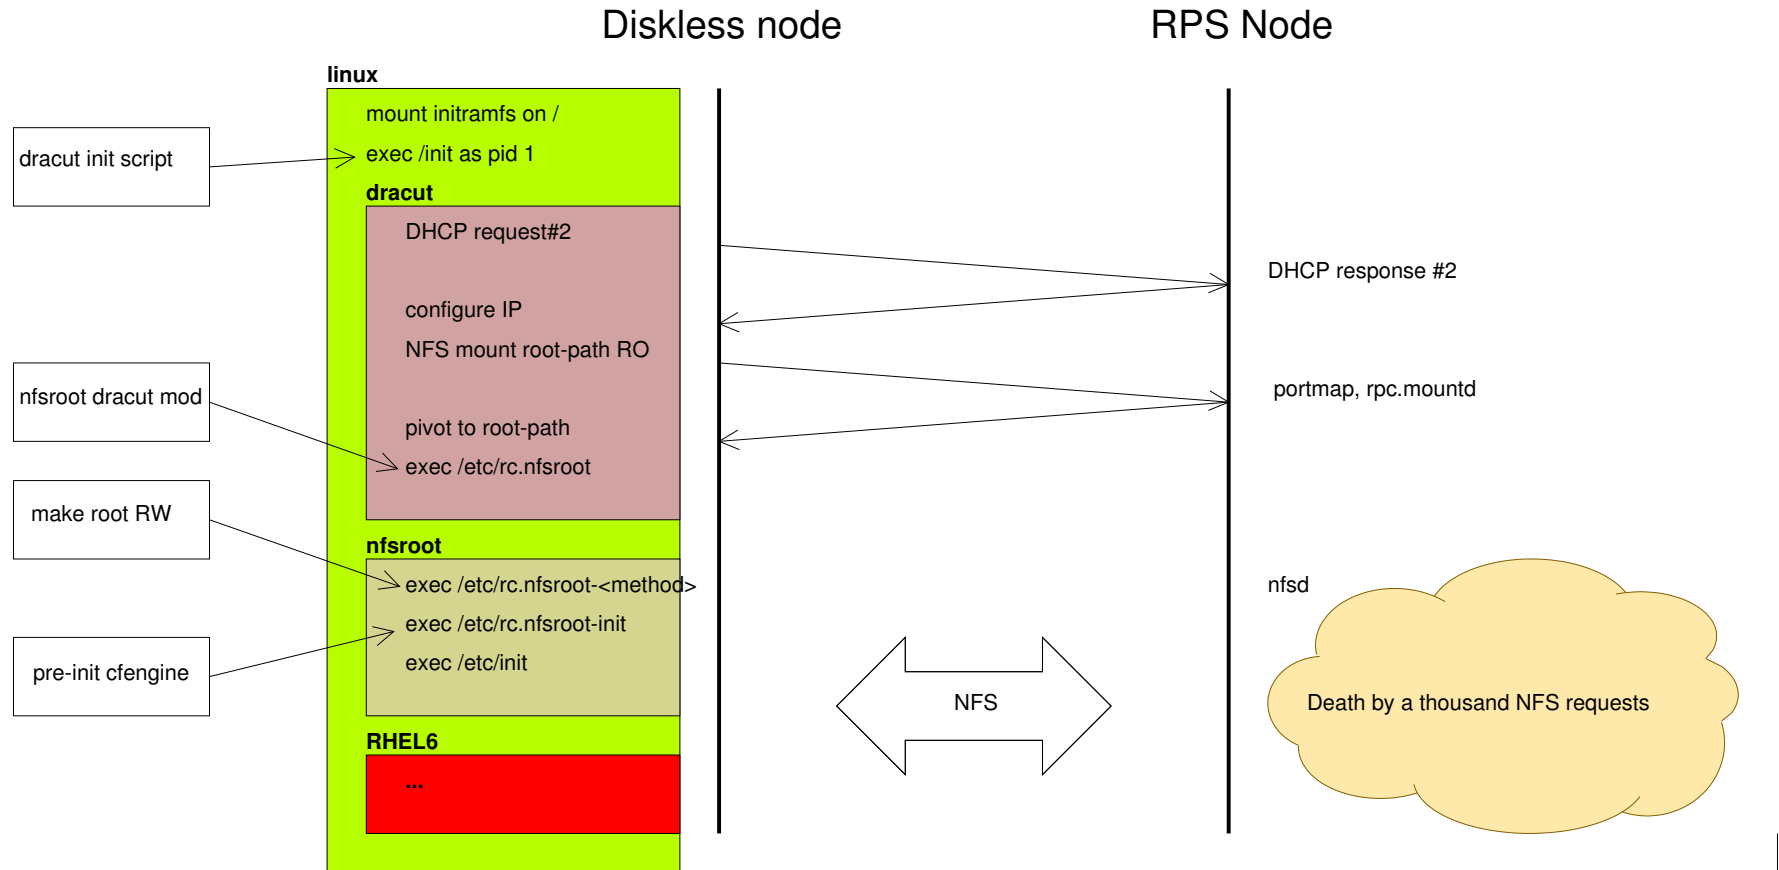
bootstrap - BIOS

Diskless node

RPS Node




bootstrap - Linux



scalability questions for large clusters

current scalability picture



- 5m to boot zin (18 SU * 162 = 2916 nodes)
- smaller clusters boot in about the same time
- scalable unit strategy seems to be working pretty well, but
- zin ARP table had to be hardwired
- bcasts travel widely
- RPS mirror rsync
- NFS caching not ideal

possible scalability improvements

- Multicast TFTP (RFC 2090)?
TFTP only transfers about $16\text{MB} * 161 = 2.5\text{GB}$ per SU
- boot-over-Infiniband?
Requires IB card boot agent, but b/w not really a problem.
- IPv6 Neighbor Discovery Protocol (NDP)?
May solve ARP scalability without hardwiring.
- IP subnetting
- NFSv4 delegations?
Reduce revalidation traffic when root image is static?
- distributed network block device?
Enables aggressive caching and distribution.

distributed network block device

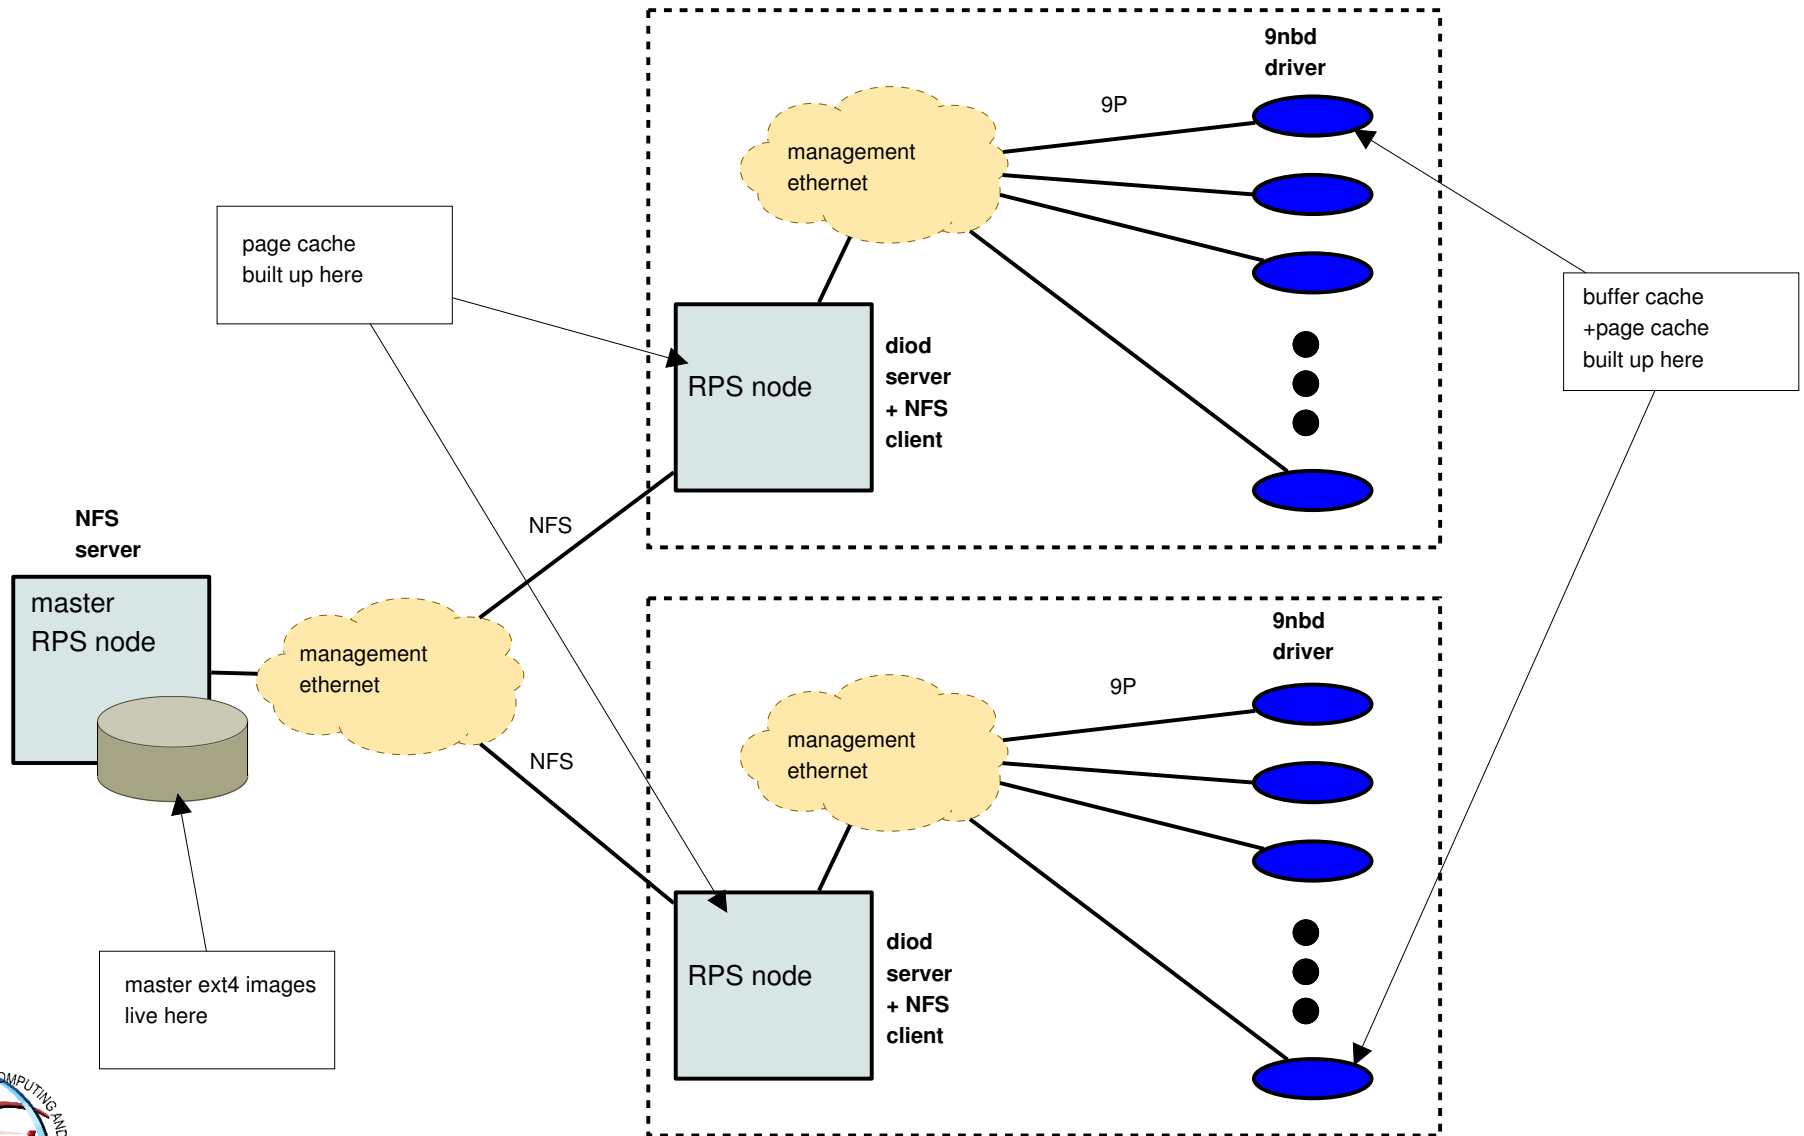
path walk: NFS scalability

- NFS does not scale well for PATH search, LD_LIBRARY search, python module search, etc
- NFS doesn't cache whole directories.
- dcache (+/-) only helps for names looked up before
- round-trip to NFS server for each unknown lookup
- NFS timeout based revalidation

path walk: block device scalability

- dracut supports nbd, iscsi root
- root uses a local file system (ext4, squashfs, ...)
- nfsroot **zram** method added in 3.20
RO block device + RW zram device using lvm-snapshot
- blocks backing directories are cached
- no revalidation
- caveat: RO root image must not change underneath
diskless nodes!

distributed network block device



pathwalk test: NFS vs 9NBD

Search for 10K names in 16 dirs, each containing 10K files.

